

Chapter 4 — Best practices for annotation

Use cases illustrate stakes; day-to-day habits sustain quality

Learning objectives

- Draft guideline habits with positive and negative exemplars
- Use pre-labeling only with audit thresholds met

Guidelines and exemplars

- Strong annotation programs share a common skeleton
- Guidelines specify class definitions, edge cases, and prohibited inferences
- Include positive and negative examples for each label and document how to handle ambiguity
- Version the guideline document when schemas change, retrain annotators on diffs

Feedback, audits, and drift checks

- Run regular calibration sessions while labeling is active
- Insert repeated gold items to detect drift over long projects

Diversity and bias reduction

- Where labels encode social constructs such as sentiment
- In practice: include diverse reviewers, test subgroup agreement
- Diversity is not only a collection concern

Pre-labeling and effort reduction

- Use pre-labeling to propose candidates
- Combine active learning, pre-labeling
- A model proposal accepted without correction is still a human decision and should appear

Takeaways

- Versioned guidelines with inline exemplars anchor every labeler
- Calibration, gold items, and re-audits catch drift after tool or schema changes
- Pre-labeling saves time only when proposals stay under human review with measured error

Next

- Complete the quiz for this part
- The bridge to data cleaning in Chapter 5